

Introduction to Information Retrieval (Chapter 6)

Scoring, term weighting, and the vector space model)

As for Boolean queries, the resulting number of matching documents can far exceed the number a human user could possibly sift through. Consequently, it is essential for a search engine to rank-order the documents matching a query. To rank the search results, we have to provide the scores of results.

1. Parametric and zone indexes

- **Parametric**

Query: *find documents authored by William Shakespeare in 1601, containing the phrase alas poor Yorick.*

There is one parametric index for each field (date of creation). It allows us to select only the documents matching a date specified in the query.

- **Zones**

Query: *find documents with merchant in the title and william in the author list and the phrase gentle rain in the body.*

Zones are similar to fields, except the contents of a zone can be arbitrary free text. For instance, document titles and abstracts are generally treated as zones.

- **Weighted zone scoring**

Give a Boolean query q and a document d , where each zone of the document contributes a Boolean value. Weighted zone score is a linear combination of zone scores.

Let s_i be the Boolean score denoting a match (or absence thereof) between q and the i th zone. Let $g_1, \dots, g_l \in [0, 1]$ such that $\sum_{i=1}^l g_i = 1$. Then the weighted zone score is defined to be:

$$\sum_{i=1}^l g_i s_i \in [0, 1] \quad (1)$$

The algorithm in Figure 6.4 treats the case when the query q is a two-term query consisting of query terms q_1 and q_2 , and the Boolean function is AND: 1 if both query terms are present in a zone and 0 otherwise.

ZONE SCORE(q_1, q_2)

```

1  float scores[N] = [0]
2  constant g[l]
3   $p_1 \leftarrow \text{postings}(q_1)$ 
4   $p_2 \leftarrow \text{postings}(q_2)$ 
5  // scores[] is an array with a score entry for each document, initialized to zero.
6  //  $p_1$  and  $p_2$  are initialized to point to the beginning of their respective postings.
7  // Assume g[] is initialized to the respective zone weights.
8  while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$ 
9  do if  $\text{docID}(p_1) = \text{docID}(p_2)$ 
10     then scores[docID( $p_1$ )]  $\leftarrow \text{WEIGHTEDZONE}(p_1, p_2, g)$ 
11          $p_1 \leftarrow \text{next}(p_1)$ 
12          $p_2 \leftarrow \text{next}(p_2)$ 
13     else if  $\text{docID}(p_1) < \text{docID}(p_2)$ 
14         then  $p_1 \leftarrow \text{next}(p_1)$ 
15         else  $p_2 \leftarrow \text{next}(p_2)$ 
16  return scores

```

Figure 6.4 Algorithm for computing the weighted zone score from two postings lists. Function WEIGHTEDZONE (not shown here) is assumed to compute the inner loop of Equation 6.1.

2. Tf-idf function

2.1 Tf-idf

- **Term frequency (TF):** the number of occurrences of term t in document d , note as $tf_{t,d}$.
- **Inverse document frequency (idf):** $idf_t = \log \frac{N}{df_t}$, where N is the total number of documents in a collection; df_t , document frequency, is the number of documents in the collection that contain a term t .
- **TF-idf:** $tf-idf_{t,d}$ assigns to term t a weight in document d :

$$tf-idf_{t,d} = tf_{t,d} \times idf_t \quad (2)$$

2.2 Variant tf-idf functions

2.2.1 Sublinear tf scaling

According to the definition of tf , the importance of a term is proportional to its occurrence in the document, the twenty occurrences of a term in a document truly carry twenty times the significance of a single occurrence. However, this conclusion may not be correct. A common modification is to use instead of the logarithm of the term frequency, which assigns a weight given by

$$w_{tf,d} = \begin{cases} 1 + \log tf_{t,d} & \text{if } tf_{t,d} > 0 \\ 0 & \text{otherwise} \end{cases} \quad (3)$$

2.2.2 Maximum tf normalization

Suppose we were to take a document d and create a new document d' by simply appending a copy of d to itself. Although d' should be no more relevant to any query than d is, the use of tf-idf assigns it twice as high a score as d . Therefore, replacing tf-idf by ntf-idf eliminates the anomaly in this example.

The main idea of maximum tf normalization is to mitigate the following anomaly: We observe higher term frequencies in longer documents, merely because longer documents tend to repeat the same words over and over again.

Let $tf_{max}(d) = \max_{\tau \in d} tf_{\tau,d}$, where τ ranges over all terms in d . A normalized term frequency for each term τ in document d by

$$ntf_{t,d} = a + (1 - a) \frac{tf_{t,d}}{tf_{max}(d)} \quad (4)$$

where a is a value between 0 and 1 and is generally set to 0.4, although some early work used the value 0.5. The term a in (4) is a smoothing term whose role is to damp the contribution of the second term, which may be viewed as a scaling down of tf by the largest tf value in d .

2.2.3 Document and query weighting schemes

The document and query vector can be represented by the SMART notation:

term frequency	document frequency	normalization
n (natural) $tf_{t,d}$	n (no) 1	n (none) 1
l (logarithm) $1 + \log(tf_{t,d})$	t (idf) $\log \frac{N}{df_t}$	c (cosine) $\frac{1}{\sqrt{w_1^2 + w_2^2 + \dots + w_M^2}}$
a (augmented) $0.5 + \frac{0.5 \times tf_{t,d}}{\max_t(tf_{t,d})}$	p (prob idf) $\max\{0, \log \frac{N - df_t}{df_t}\}$	u (pivoted unique) $1/u$ (Section 17.4.4)
b (boolean) $\begin{cases} 1 & \text{if } tf_{t,d} > 0 \\ 0 & \text{otherwise} \end{cases}$		b (byte size) $1/CharLength^\alpha$, $\alpha < 1$
L (log ave) $\frac{1 + \log(tf_{t,d})}{1 + \log(\text{ave}_{t \in d}(tf_{t,d}))}$		

Figure 6.15 SMART notation for tf-idf variants. Here *Char Length* is the number of characters in the document.

We could combine different document and query weighting schemes (term frequency - document frequency - normalization) based on SMART. A very standard weighting scheme is Inc.Itc, where the document vector has log-weighted term frequency, no idf (for both effectiveness and efficiency reasons), and cosine normalization, while the query vector uses log-weighted term frequency, idf weighting, and cosine normalization.

Summary

- A very standard weighting scheme is Inc.Itc, where the document vector has log-weighted term frequency, no idf (for both effectiveness and efficiency reasons), and cosine normalization, while the query vector uses log-weighted term frequency, idf weighting, and cosine normalization.