

# Introduction to Information Retrieval (Chapter 11)

## Probabilistic information retrieval

### 1. The binary independence model

- “Binary” is equivalent to Boolean: Documents and queries are both represented as binary term incidence vectors. That is, a document  $d$  is represented by the vector  $\vec{x} = (x_1, \dots, x_M)$  where  $x_t = 1$  if term  $t$  is present in document  $d$  and  $x_t = 0$  if  $t$  is not present in  $d$ . Similarly, we represent  $q$  by the incidence vector  $\vec{q}$ .
- “Independence” means that terms are modeled as occurring in documents independently.

Based on Bayes rule, we have:

$$\begin{aligned} P(R = 1 | \vec{x}, \vec{q}) &= \frac{P(\vec{x} | R = 1, \vec{q}) P(R = 1 | \vec{q})}{P(\vec{x} | \vec{q})} \\ P(R = 0 | \vec{x}, \vec{q}) &= \frac{P(\vec{x} | R = 0, \vec{q}) P(R = 0 | \vec{q})}{P(\vec{x} | \vec{q})} \end{aligned} \quad (1)$$

Here,  $P(R = 1 | \vec{x}, \vec{q})$  and  $P(R = 0 | \vec{x}, \vec{q})$  are the probability that if a relevant or nonrelevant, respectively, document is retrieved, then that document's representation is  $\vec{x}$ .  $P(R = 1 | \vec{q})$  and  $P(R = 0 | \vec{q})$  indicate the prior probability of retrieving a relevant or nonrelevant document, respectively, for a query  $\vec{q}$ .

Odds, a kind of multiplier for how probabilities change, is monotonic with the probability of relevance:

$$\text{odds } O(A) = \frac{P(A)}{P(\bar{A})} = \frac{P(A)}{1 - P(A)} \quad (2)$$

Thus, the odds of relevance is:

$$O(R | \vec{x}, \vec{q}) = \frac{P(R = 1 | \vec{x}, \vec{q})}{P(R = 0 | \vec{x}, \vec{q})} = \frac{P(R = 1 | \vec{x}, \vec{q})}{P(R = 1 | \vec{x}, \vec{q})} \cdot \frac{P(R = 1 | \vec{q})}{P(R = 0 | \vec{q})} \cdot \frac{P(\vec{x} | R = 1, \vec{q})}{P(\vec{x} | R = 0, \vec{q})} \quad (3)$$

According to Naive Bayes conditional independence assumption,

$$\frac{P(\vec{x} | R = 1, \vec{q})}{P(\vec{x} | R = 0, \vec{q})} = \prod_{t=1}^M \frac{P(x_t | R = 1, \vec{q})}{P(x_t | R = 0, \vec{q})} \quad (4)$$

So,

$$O(R | \vec{x}, \vec{q}) = O(R | \vec{q}) \cdot \prod_{t=1}^M \frac{P(x_t | R = 1, \vec{q})}{P(x_t | R = 0, \vec{q})} \quad (5)$$

Since each  $x_t$  is either 0 or 1, we can separate the terms to give:

$$O(R | \vec{x}, \vec{q}) = O(R | \vec{q}) \cdot \prod_{t=1}^M \frac{P(x_t = 1 | R = 1, \vec{q})}{P(x_t = 1 | R = 0, \vec{q})} \cdot \prod_{t=1}^M \frac{P(x_t = 0 | R = 1, \vec{q})}{P(x_t = 0 | R = 0, \vec{q})} \quad (6)$$

Henceforth, let  $p_t = P(x_t = 1 | R = 1, \vec{q})$  be the probability of a term appearing in a document relevant to the query, and  $u_t = P(x_t = 1 | R = 0, \vec{q})$  be the probability of a term appearing in a nonrelevant document. These quantities can be visualized in the following contingency table where the columns add to 1:

document		relevant ( $R = 1$ )	nonrelevant ( $R = 0$ )
term present	$x_t = 1$	$p_t$	$u_t$
term absent	$x_t = 0$	$1 - p_t$	$1 - u_t$

table 1.

Considering if terms are not occurring in the query, there is insignificant to calculate  $p_t$  and  $u_t$ . Thereby, we only consider terms in the products that appear in the query ( $q_t = 1$ ), so

$$O(R|\vec{x}, \vec{q}) = O(R|\vec{q}) \cdot \prod_{t:x_t=q_t=1} \frac{p_t}{u_t} \cdot \prod_{t:x_t=0, q_t=1} \frac{1-p_t}{1-u_t} \quad (7)$$

Because,

$$\prod_{t:x_t=0, q_t=1} \frac{1-p_t}{1-u_t} = \prod_{t:x_t=0, q_t=1} \frac{1-p_t}{1-u_t} \cdot \prod_{t:x_t=1, q_t=1} \frac{1-p_t}{1-u_t} \frac{1-u_t}{1-p_t} = \prod_{t:q_t=1} \frac{1-p_t}{1-u_t} \cdot \prod_{t:x_t=1, q_t=1} \frac{1-u_t}{1-p_t}$$

Therefore,

$$O(R|\vec{x}, \vec{q}) = O(R|\vec{q}) \cdot \prod_{t:x_t=q_t=1} \frac{p_t(1-u_t)}{u_t(1-p_t)} \cdot \prod_{t:q_t=1} \frac{1-p_t}{1-u_t} \quad (8)$$

Note,

$$RSV_d = \log \prod_{t:x_t=q_t=1} \frac{p_t(1-u_t)}{u_t(1-p_t)} = \sum_{t:x_t=q_t=1} \log \frac{p_t(1-u_t)}{u_t(1-p_t)} \quad (9)$$

$$c_t = \log \frac{p_t(1-u_t)}{u_t(1-p_t)} = \log \frac{p_t}{(1-p_t)} + \log \frac{1-u_t}{u_t} \quad (10)$$

where, the resulting quantity used for ranking is retrieval called the retrieval status value (RSV), the  $c_t$  terms are log odds ratios for the terms in the query.

Consequently,

$$RSV_d = \sum_{t:x_t=q_t=1} c_t \quad (11)$$

Based on table 1, count the number of documents,

documents		relevant	nonrelevant	total
term present	$x_t = 1$	$s$	$df_t - s$	$df_t$
term absent	$x_t = 0$	$S - s$	$(N - df_t) - (S - s)$	$N - df_t$
total		$S$	$N - S$	$N$

table 2.

In table 2,  $df_t$  is the number of documents that contain term t, and  $p_t = s/S$ ,  $u_t = (df_t - s)/(N - S)$ , so,

$$c_t = K(N, df_t, S, s) = \log \frac{s/(S-s)}{(df_t - s)/((N - df_t) - (S - s))} \quad (12)$$

To avoid the possibility of zeroes (such as if every or no relevant document has a particular term) it is fairly standard to add  $1/2$  to each of the quantities in the center four terms of (12), and then to adjust the marginal counts (the totals) accordingly (so, the bottom right cell totals  $N + 2$ ). Then we have:

$$c_t = K(N, df_t, S, s) = \log \frac{(s + \frac{1}{2}) / (S - s + \frac{1}{2})}{(df_t - s + \frac{1}{2}) / ((N - df_t) - (S - s + \frac{1}{2}))} \quad (13)$$

## 2. An appraisal and some extensions

Getting reasonable approximations of the needed probabilities for a probabilistic IR model is possible, but it requires some major assumptions. In the BIM these are:

- a Boolean representation of documents/queries/relevance
- term independence
- terms not in the query don't affect the outcome
- document relevance values are independent

### 2.1 Okapi BM25: A nonbinary model

The BIM was originally designed for short catalog records and abstracts of fairly consistent length, and it works reasonably in these contexts, but for modern full-text search collections, it seems clear that a model should pay attention to term frequency and document length.

The BM25 weighting scheme, often called Okapi weighting, after the system in which it was first implemented, was developed as a way of building a probabilistic model sensitive to these quantities while not introducing too many additional parameters into the model (Sparck Jones et al. 2000).

The simplest score for document  $d$  is just idf weighting of the query terms present,

$$RSV_d = \sum_{t \in q} \log \frac{N}{df_t} \quad (14)$$

We can improve on Equation (14) by factoring in the frequency of each term and document length:

$$RSV_d = \sum_{t \in q} \log \left[ \frac{N}{df_t} \right] \cdot \frac{(k_1 + 1)tf_{td}}{k_1((1 - b) + b \times (L_d/L_{ave})) + tf_{td}} \quad (15)$$

Here,  $tf_{td}$  is the frequency of term  $t$  in document  $d$ , and  $L_d$  and  $L_{ave}$  are the length of document  $d$  and the average document length for the whole collection. The variable  $k_1$  is a positive tuning parameter that calibrates the document term frequency scaling. A  $k_1$  value of 0 corresponds to a binary model (no term frequency), and a large value corresponds to using raw term frequency.  $b$  is another tuning parameter ( $0 \leq b \leq 1$ ) that determines the scaling by document length:  $b = 1$  corresponds to fully scaling the term weight by the document length, whereas  $b = 0$  corresponds to no length normalization.

If the query is long, then we might also use similar weighting for query terms. This is appropriate if the queries are paragraph-long information needs, but unnecessary for short queries.

$$RSV_d = \sum_{t \in q} \log \left[ \frac{N}{df_t} \right] \cdot \frac{(k_1 + 1)tf_{td}}{k_1((1 - b) + b \times (L_d/L_{ave})) + tf_{td}} \cdot \frac{(k_3 + 1)tf_{tq}}{k_3 + tf + tq} \quad (15)$$

with  $tf_{tq}$  being the frequency of term  $t$  in the query  $q$ , and  $k_3$  being another positive tuning parameter that this time calibrates term frequency scaling of the query. In the absence of such optimization, experiments have shown reasonable values are to set  $k_1$  and  $k_3$  to a value between 1.2 and 2 and  $b = 0.75$ .

The BM25 term weighting formulas have been used quite widely and quite successfully across a range of collections and search tasks. Especially in the TREC evaluations, they performed well and were widely adopted by many groups.

## 2.2 Bayesian network approaches to information retrieval

Probabilistic network is the primary case of a statistical ranked retrieval model that naturally supports structured query operators. The system allowed efficient large-scale retrieval, and was the basis of the InQuery text retrieval system, built at the University of Massachusetts. This system performed very well in TREC evaluations and for a time was sold commercially.

## Conclusions

- The difference between “vector space” and “probabilistic” IR systems is not that great. For a probabilistic IR system, it’s just that, at the end, you score queries not by cosine similarity and tf–idf in a vector space, but by a slightly different formula motivated by probability theory. Indeed, sometimes people have changed an existing vector-space IR system into an effectively probabilistic system simply by adopted term weighting formulas from probabilistic models.
- The probabilistic approach to IR originated in the United Kingdom in the 1950s. The first major presentation of a probabilistic model is Maron and Kuhns (1960). Robertson and Jones (1976) introduce the main foundations of the BIM and van Rijsbergen (1979) presents in detail the classic BIM probabilistic model.